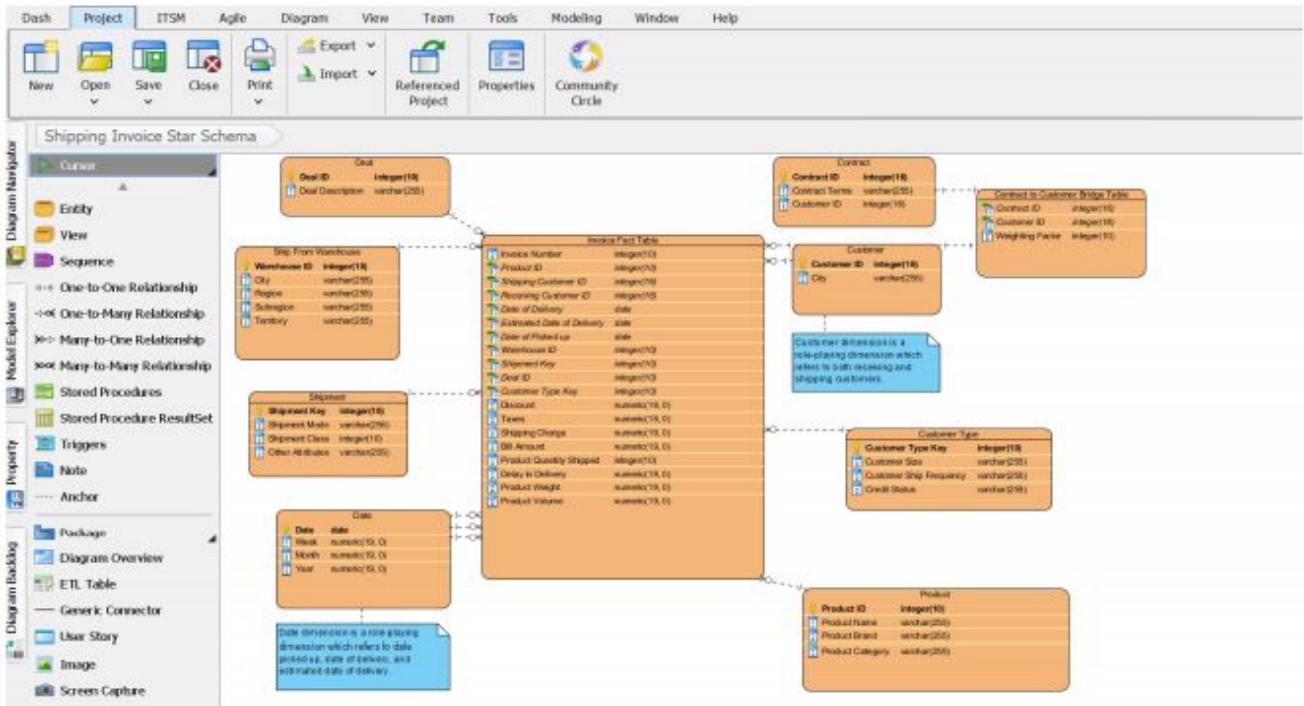


## Part A: Star Schema Exercise - Data Warehouse Modeling

Important notes:

- The following diagram is designed using Visual Paradigm.
- All descriptions are selectively quoted from the 543 Final instructions. I have asked my professor for approval.



The descriptions of the Star Schema are:

- A shipping company serves a limited group of customers.
- A shipment invoice contains the following information:
  - Invoice Number; Date Picked Up; Date of Delivery; Ship From Customer; Ship To Customer; Ship From Warehouse; Ship From City; Ship To City; Shipment Mode (like air/sea/truck/train); Shipment Class (codes like 1,2,3 which translate to express, expedited, standard); Contract ID; Total Shipping Charges; Deal ID (refers to a discount deal in effect); Discount Amount; Taxes; Total Billed Amount; Total Ship Weight; Total Ship Volume
- Each Line in Invoice contains: Product ID, Product Name, Quantity, Ship Weight, Ship Volume, Charge

Assumptions:

- Quantity is number of items
- Discount and Taxes can be allocated per line
- Number of contracts is relatively small. Each customer may have multiple contracts
- Number of Deals is relatively small. Some shipments may have no deal applicable

Requirement:

- Design a dimensional model, with invoice line as the grain of the fact table.

Fact Table (shipment invoice line)

Invoice Number (PK, DD)

Invoice\_Product (FK)

Deal ID (FK)

Date Picked Up Key (FK)

Date of Delivery Key (FK)

*[both access Date Dimension]*

\*Customer ship from Key (FK)

Ship from Key (FK) - inside this dim table = correlated warehouse and city info

Shipment\_Info (FK) - inside this table = correlated modes and class info

Ship to Key (FK)

\*Customer ship to Key (FK)

*[both access Customer Dimension]*

Contract ID (FK)

Discount Dollar Amount

Taxes Dollar Amount

Total Ship Weight

Total Ship Volume

Total Billed Amount

*[Additional attributes not necessary in this question but in real life may include:]*

Customer satisfaction key (FK)

Fixed cost

Variable cost

Distribution cost

Contribution dollar amount

Shipment line item on time count

Shipment line item complete count

Shipment line item damage free count

Dimension Tables:

Date Dimension: Key, Date, other attributes

Customer Dimension: cust\_key, name, street address, city, state, zip, phone, other attributes

*Potential outlier table:* attributes for customer size, customer ship frequency, credit status (low cardinality, depends on how often this information is queried)

Invoice\_Product Dimension: (invoice number, product ID), product name, category, quantity, individual product ship weight, individual product ship volume, individual product charge, other attributes  
(1 to many relationship with fact table)

Ship-from Dimension: key, warehouse name, street, city, state, contact person, facility type, subregion, region, territory, other attributes  
*Contains one row for each manufacturer warehouse or shipping location*

Ship-to Dimension: key, warehouse name, street, city, state, contact person, facility type, subregion, region, territory, other attributes  
*Could consider merging ship from and ship to tables (i.e. larger ship\_id table, perhaps w. outtrigger for storing overlapping information...depends on pragmatics of querying)*

Shipment\_Info Dimension: ID, shipment mode, shipment class, total weight, total volume, total charges, other attributes  
Descriptions for a particular shipment mode + class for estimated delivery duration

Contract Dimension: ID, contract duration start/end, Contract Terms, Policies, handling errors/delays in delivery, other attributes  
*Assuming that each customer may have multiple contracts, but a unique contract would only apply to 1 customer*

Deal Dimension: ID, deal terms, discount rates, other attributes

**Question:** At the moment we cannot track information about deals in existence on any given date with the fact table of invoice line items. We would need to join the FT and deal dimensional table (since only the deal ID is available immediately in the FT). When we need detailed attributes like existence by date, we would need to query something like:

```
SELECT interested_deal_attributes
FROM Deal
INNER JOIN Invoice
ON Deal.Deal_ID = Invoice.Deal_ID
WHERE Date_of_Delivery = givenDate
```

Another option might be creating a mini-dimension with just this required information (subset of attributes from deal dimension table) for faster querying, depending on business pragmatics.

## Part B:

1. State what each of the following is **briefly**. Provide an example for each.

a. MiniDimension

Stores a subset of attributes which change frequently (or are often queried together in combination) from a larger dimension, and becomes connected to the fact table as a separate dimension. Minidimensions consist of correlated clumps of attributes and support consistent day-to-day analytic reporting (business pragmatics).

Example: A large customer dimension includes attributes including age, number of children, income level. If these particular attributes are changed or analyzed frequently, we could have them removed to create a demographics minidimension (with a new FK added to the fact table) for better efficiency. Attributes not as frequently queried are left in the original large cust. Dim.

b. Junk Dimension

Stores miscellaneous attributes (i.e. low cardinality indicators, flag codes, text) that are unrelated to any particular dimension, and hence stored in a separate table for convenience. These attributes are not accessed frequently.

Example: A hospital's DW has a gender dimension and marital status dimension. In the fact table, there are currently 2 keys referring to these dimensions (if kept as single dimensions, would be limited to single attributes). To eliminate these small dimensions, a junk dimension cross joins all possible attributes (gender, marital status) into a single dimension, with just a single key in the fact table.

c. Degenerate Dimension

A dimension key (in the fact table) that does not join to a corresponding dimension table, and has no attributes. Common when the grain of a FT represents a single transaction item/line item

Example: Operational control #s: Order numbers, invoice numbers, bill-of-lading numbers  
These can be maintained in the fact table instead of in separate dimensions, as we have already extracted info (attributes) from other dimensions. A more specific example: PK of the retail sales FT consists of a degenerate POS transaction number and product key (assuming the POS system rolls up all sales for a given product within a POS shopping cart into a single line item).

d. MultiValued Dimension

A dimension with more than 1 value per fact row, compared to the classic dimensional schema where each dimension attached to the FT has a single value. Bridge tables are often used.

Example: A patient receiving a healthcare treatment may have multiple concurrent diagnoses. The multivalued dimension is attached to the FT via a group dimension key to a bridge table, with one row for each concurrent diagnosis in a group.

Another example (textbook): Bank account can have one, two, or more individual account holders. Including customer as an account attribute violates the granularity of the dimension table. And including customer as an additional dimension in the fact table violates the granularity of the FT, since more than one individual can be associated with any given account.

*2. You created an Employee dimension for a company with 50,000 employees, and identified Team Assignment as an attribute that changes every six months. Management wants to track history of Team assignment. How will you handle change of this attribute?*

I would use SCD type 2 to track history of Team Assignment, that is, inserting new rows to capture the assignment changes (using surrogate keys). This technique is powerful because the new dimension row automatically partitions history in the fact table, and is gracefully extensible (not a re-design of the tables). I believe this method is appropriate for now, because Kimball & Ross mentions type 2 becomes problematic when dimension tables already exceed a million rows. Since we are basically adding 50,000 rows every 6 months, we can use this method to store 10 years' worth of data before reaching this threshold.

*In another project, you find a similar situation, with team assignment changing every two weeks in a very large employee dimension. Management wants to track history of team assignment here as well. How will you handle change of this attribute?*

After initially considering hybrid approaches, I realized that with the fast (and large) changes happening every 2 weeks, we're not really dealing with 'slowly' changing dimension anymore (so the previous SCD techniques aren't optimal...it's a fast changing dimension). My approach would be breaking off the 'Team Assignment' attribute from the Employee dimension (along with any other attributes changed at this frequency) into a separate minidimension. The fact table would then have two foreign keys—one for the primary Employee dimension table and another for the rapidly changing 'Team Assignment' attribute. These dimension tables would be associated with one another every time we put a row in the fact table.

3. A. Give an example of what could be a mini-dimension that spins off from one of the dimensions in this diagram (make suitable assumptions). Identify which dimension you take the mini-dimension out from, what kind of information may be part of the mini-dimension, and the reasoning behind spinning off the mini-dimension.

One example would be 'Policyholder Dimension', as we may want to make mini-dimension(s) containing attributes like policyholder's ZIP code (also for number of children, income level, etc.). The ZIP code would change as the policyholder moves, and is closely correlated to the insurer's pricing and risk algorithms for these customers. The company wants to track these changes, instead of just overwriting records. We may assume that a national insurance company will have well over 1 million rows in the large policyholder dimension table. If true, there would be significant burden on the data staging (as the number of rows grow) to support tracking of these fast changing attributes. So splitting into a minidimension directly linked to the fact table with a separate surrogate key will improve the efficiency of attribute browsing (since the company needs to frequently query + update these attributes). Whether ZIP code is the only attribute in a mini-dimension (versus combining other attributes like income level) depends on how frequently these attributes are queried together (correlation/business pragmatics), and whether updating occurs at similar periods. Other examples may be found in the covered item dimension.

B. There are different kinds of employees who handle the different transaction types, for example, these could be sales rep, agent, auditor. Currently they are all part of the single Employee dimension. Let us say that individual employee types have some attributes unique to each, while they continue to share the common attributes currently in the Employee dimension. If we also wanted to include the custom employee type attributes into the schema, how will we do that?

Instinctively I think this would be a good example of a permissible outrigger, that is, creating another table (connected to the Employee dimension). If this question implies the common attributes should continue to stay in the Employee dimension, then we'd place the custom employee type attributes into the outrigger (for more detailed querying). This is assuming that these custom employee type attributes have low-cardinality and perhaps aren't queried super frequently. (If the implication is false, we might switch the contents of the outrigger and Employee dimension tables around, depending on what reduces redundancy of repeated information).

But these decisions are ultimately driven by business pragmatics. If somehow these custom employee type attributes become super frequently analyzed in the future, we might instead consider placing them into a mini-dimension connected to the FT for operationally faster queries.